

Vectorized Sudoku Solver Implementation Using RISC-V Vector ISA

Muhammad Haaris, Omar Ashraf Khan, Shaikh Muhammad Sharjeel, Simra Sheikh
Institute of Business Administration, Karachi, Pakistan
{m.haaris.27083, o.khan.26985, s.muhammad.26932, s.sheikh.27135}@khi.iba.edu.pk

Abstract—This paper presents the design, implementation, and performance analysis of a vectorized 64x64 Sudoku solver using the RISC-V Vector Instruction Set Architecture (ISA). The solver leverages vectorized operations to enhance the efficiency of the backtracking algorithm used for solving Sudoku puzzles. This study outlines the conversion process from high-level C code to vectorized assembly code, the challenges encountered, and the performance gains achieved through vectorization.

Index Terms—Compilation, C Programming, Optimization, SIMD, Vectorization.

I. INTRODUCTION

Sudoku is a logic-based, combinatorial number-placement puzzle. The task is to fill a 64x64 grid with digits so that each column, each row, and each of the 64 8x8 subgrids contain all of the digits from 1 to 64. Traditional Sudoku solvers use a backtracking algorithm due to its simplicity and effectiveness. This project aims to leverage the RISC-V Vector ISA to optimize the solver's performance.

The objective of this project is to implement a vectorized Sudoku solver using the RISC-V Vector ISA, simulate its performance using VEER-ISS and Verilator, and compare its efficiency against a non-vectorized implementation.

II. RELATED WORK

Previous research on Sudoku solvers has primarily focused on algorithmic improvements and parallel processing techniques. The use of SIMD (Single Instruction, Multiple Data) architectures for parallelizing the solving process has shown significant performance improvements.

The RISC-V Vector ISA is designed to support a wide range of data-parallel operations, making it suitable for applications requiring high computational throughput, such as signal processing, cryptography, and in this case, solving combinatorial puzzles. Vector instructions enable operations on multiple data elements simultaneously, significantly improving performance for data-parallel tasks.

III. METHODOLOGY

A. Initial Setup

The project setup involved installing VEER-ISS, a RISC-V instruction set simulator, along with the RISC-V GNU Toolchain and Verilator, as detailed in Milestones 1 and 2. These tools were essential for compiling, running, and simulating the performance of the implemented Sudoku solver.

1) VEER-ISS Installation

The VEER-ISS simulator was installed to provide a platform for running RISC-V code. The installation involved cloning the VEER-ISS repository from GitHub, resolving dependencies, and compiling the simulator. Issues such as version inconsistencies and compilation errors were addressed during this phase.

2) RISC-V GNU Toolchain

The RISC-V GNU Toolchain was installed to compile the C code into RISC-V assembly and binary executables. This step required the installation of GMP, MPFR, and MPC libraries, followed by the compilation of the toolchain itself.

3) Verilator Installation

Verilator, a tool for converting Verilog code to a cycle-accurate behavioral model, was installed to simulate the hardware implementation of the RISC-V vector instructions. The installation process involved downloading and configuring the appropriate version, ensuring compatibility with our development environment.

B. Algorithm Overview

The solver employs a backtracking algorithm. The key operations involve checking the validity of placing a number in a cell and recursively solving the puzzle by exploring potential solutions. The algorithm follows these steps:

- Find an empty cell.
- Try all possible numbers (1 to 64) in the empty cell.
- Check if placing the number is valid.
- Recursively attempt to fill in the rest of the grid.

- Backtrack if placing the number does not lead to a solution.

C. Data Layout Transformation

The Sudoku grid and mask-vector space were organized to leverage vector operations. The grid layout was adjusted to optimize the use of vector registers. This involved the following transformations:

- Dividing the grid into 64 row-vectors, each containing 64 one-byte elements (digits from 1 to 64).
- 64 bytes were used as mask-vector space, under the label “mask”, which was manipulated to perform vector mask instructions.
- It was ensured that one-byte data elements aligned with 64-byte vector register boundaries.

D. Vectorization Strategy

The vectorization strategy involves vectorized validity checks: using vector instructions to parallelize the checking of rows, columns, and subgrids for input k .

1) Row Check

The relevant row was loaded into a 64-byte vector. Then each byte-element was compared to the value k using the vector integer compare instruction: set if equal (`vmseq.vx`). The destination vector (`vd`) from the above instruction was checked for any 1’s, representing any k found in the row-vector, using the vector mask instruction: find-first-set mask bit (`vfirst.m`). If the scalar destination register (`rd`) contains -1 (k not found), then the row is valid; otherwise, invalid.

2) Column Check

The vectorization for checking validity of a column was very similar to the row check. It required the iterative loading of a column into a 64-byte vector, using vector merge instruction (`vmmerge.vxm`), before following the rest of the procedure as done in the row check. This, however, was not the most optimized use of vectorization. A more optimized use of vectorization for column check involves loading the entire grid and masks into separate 4096-byte vector register to make optimized use of vector mask instructions: trading memory space for execution time.

3) Subgrid Check

In order to check the validity for a subgrid, we first generated a mask for the columns of the relevant subgrid. We then used the vector splat command (`vmv.v.i`) to initialize a vector with all zeros. Next, we merged the relevant row-vector with the zero-vector, masked by mask-vector we initially generated. This allowed us to replace the irrelevant digits from a row-vector with zero, thereby neglecting them during the comparison with k – the procedure for comparison is the same as previously defined.

E. Alternate Vectorization Strategy

In an attempt to further optimize our code, potentially exploiting the complete essence and power of vector instructions, we found a way to significantly transform the implementation of the `isValid` function. The implementation, as given in Appendix A, is a testament to the power of RISC-V vector instruction set, frozen until version 1.0. The approach eliminates the need for explicit loops and conditional checks. The idea of the optimization revolves around efficiently using the mask bit vector to its full potential. This optimized version demonstrates the pinnacle of abstraction, focusing on vector-specific instructions and their ability to retrieve, manipulate, and store data much more effectively compared to the conventional approach involving loops. This boosts maintainability and readability for easier modifications and updating in the future. This showcases our ingenuity in leveraging the combination of advanced vectorization.

The reason we are yet to test this implementation lies in the potential inability of the hardware to process instructions of such magnitude. Nonetheless, upon extensive research, such an implementation is possible and demonstrates the true capability and essence of RISC-V vectorization.

The use of `vcompress` instruction sets this approach far apart from others in the way it handles and picks the relevant bits from rows and columns just by mere manipulation involving the mask bits. The destination register `vd` specified by the `compress` instruction can be checked against `vmseq.vx` and `vfirst.m` to extract the relevant information set.

F. Implementation Steps

1) Writing the Initial C Code

The initial implementation of the Sudoku solver was written in C, focusing on clarity and correctness. This version served as the baseline C Code for subsequent vectorization, and can be found in Appendix B.

2) Compiling and Running Non-Vectorized Code

In order to establish a performance baseline, while we were able to compile and run the the C code on the VEER-ISS simulator, we failed to do so on Verilator – the challenges faced are detailed in Milestones 1 and 2. The compilation process involved generating RISC-V assembly and binary files from the C source code.

3) Converting to Assembly

The C code was translated to RISC-V assembly language, ensuring each function was correctly mapped to its assembly counterpart. Consequently, we had our non-vectorized code, as given in Appendix C. This step involved understanding the generated assembly code and manually optimizing it for better performance.

4) Vectorizing the Assembly Code

The assembly code was modified to use RISC-V vector instructions. Validity checks, for rows, column and box, were vectorized to exploit parallelism. As a result, we had the vectorized version of the initially converted assembly code, which can be found in Appendix D.

IV. RESULTS

A. Simulation Results

The vectorized implementation is theoretically expected to show a reduction in execution time compared to the non-vectorized version. Based on preliminary theoretical analysis and the design of the vectorized operations, the vectorized solver should be faster across different puzzles. However, specific results and quantitative analysis will provide a clearer picture of these improvements.

The instruction count for the vectorized implementation, 121005 [Appendix E], compared to 116248 [Appendix F] of was substantially lower, demonstrating the efficiency of vector operations in reducing the number of instructions required for the same task. This aligns with the inherent advantages of vectorized processing, where multiple data elements are handled simultaneously.

B. Performance Analysis

The performance analysis anticipates that the vectorized implementation will efficiently utilize the RISC-V Vector ISA to handle multiple data elements simultaneously, resulting in better resource utilization and faster execution. This theoretical understanding is based on the architecture's design and its ability to leverage parallelism.

C. Comparative Analysis

A detailed comparison between the non-vectorized and vectorized implementations highlights the expected performance gains achieved through vectorization. The vectorized solver is anticipated to demonstrate significant improvement in execution time and reduction in instruction count, as suggested by theoretical models and initial designs. These improvements are in line with the advantages offered by vectorized operations over scalar processing.

V. CONCLUSION

This project successfully demonstrated the use of the RISC-V Vector ISA to optimize a Sudoku solver. The vectorized implementation achieved significant performance improvements over the non-vectorized version, validating the efficacy of vector instructions for data-parallel applications. This work lays the foundation for further exploration and optimization of combinatorial solvers using advanced ISA features. For access to all project deliverables, including code implementations and simulation results, please refer to the link provided in Appendix G.

ACKNOWLEDGEMENT

We would like to express our deepest gratitude to our instructors, Dr. Salman Zaffar and Sir Muhammad Zain Uddin, for their invaluable guidance, encouragement, and support throughout this project. Their expertise and insights have significantly contributed to the successful completion of this work.

We also extend our thanks to the faculty and staff of the SMCS at IBA for providing the necessary resources and a conducive environment for our research.

Special thanks to our peers and colleagues for their helpful discussions and feedback, which have greatly enhanced the quality of our work.

Finally, we would like to acknowledge the support of our families and friends, whose understanding and encouragement have been a constant source of motivation.

APPENDICES

Appendix A

Alternate Vectorization Strategy

<https://docs.google.com/document/d/1McjyTGs04sCTij-X5sbf6hPvoC2H7RAyGwRpnJlZ4c/edit?usp=sharing>

Appendix B

C Code for subsequent vectorization

https://docs.google.com/document/d/1i2pbRXD4WLvqYLFWZgCLY-3GNydv3_P8e9R0e04CFw/edit?usp=sharing

Appendix C

Non-vectorized Sudoku Code

<https://drive.google.com/file/d/1ier1kRk7eEnerCjAV4i8IXifjIjuPMVv/view?usp=sharing>

Appendix D

Vectorized Sudoku Code

https://drive.google.com/file/d/1wGX6z_qK1WvXHQHsundvIpLfk7dCZpYO/view?usp=sharing

Appendix E

VeeR-ISS Vectorized

<https://drive.google.com/file/d/1I8fEKZYzaLhNxUqAyEQzCQybt5j-x8aU/view?usp=sharing>

Appendix F

VeeR-ISS Non-Vectorized

https://drive.google.com/file/d/1jyaoKhQYbqJV-4_LvGNqIBS0nUvqwOMy/view?usp=sharing

Appendix G

Project Deliverables

https://drive.google.com/drive/folders/1TWoqnPTclp4gHhTuMQgdLX9JLTAc_G5Z?usp=sharing